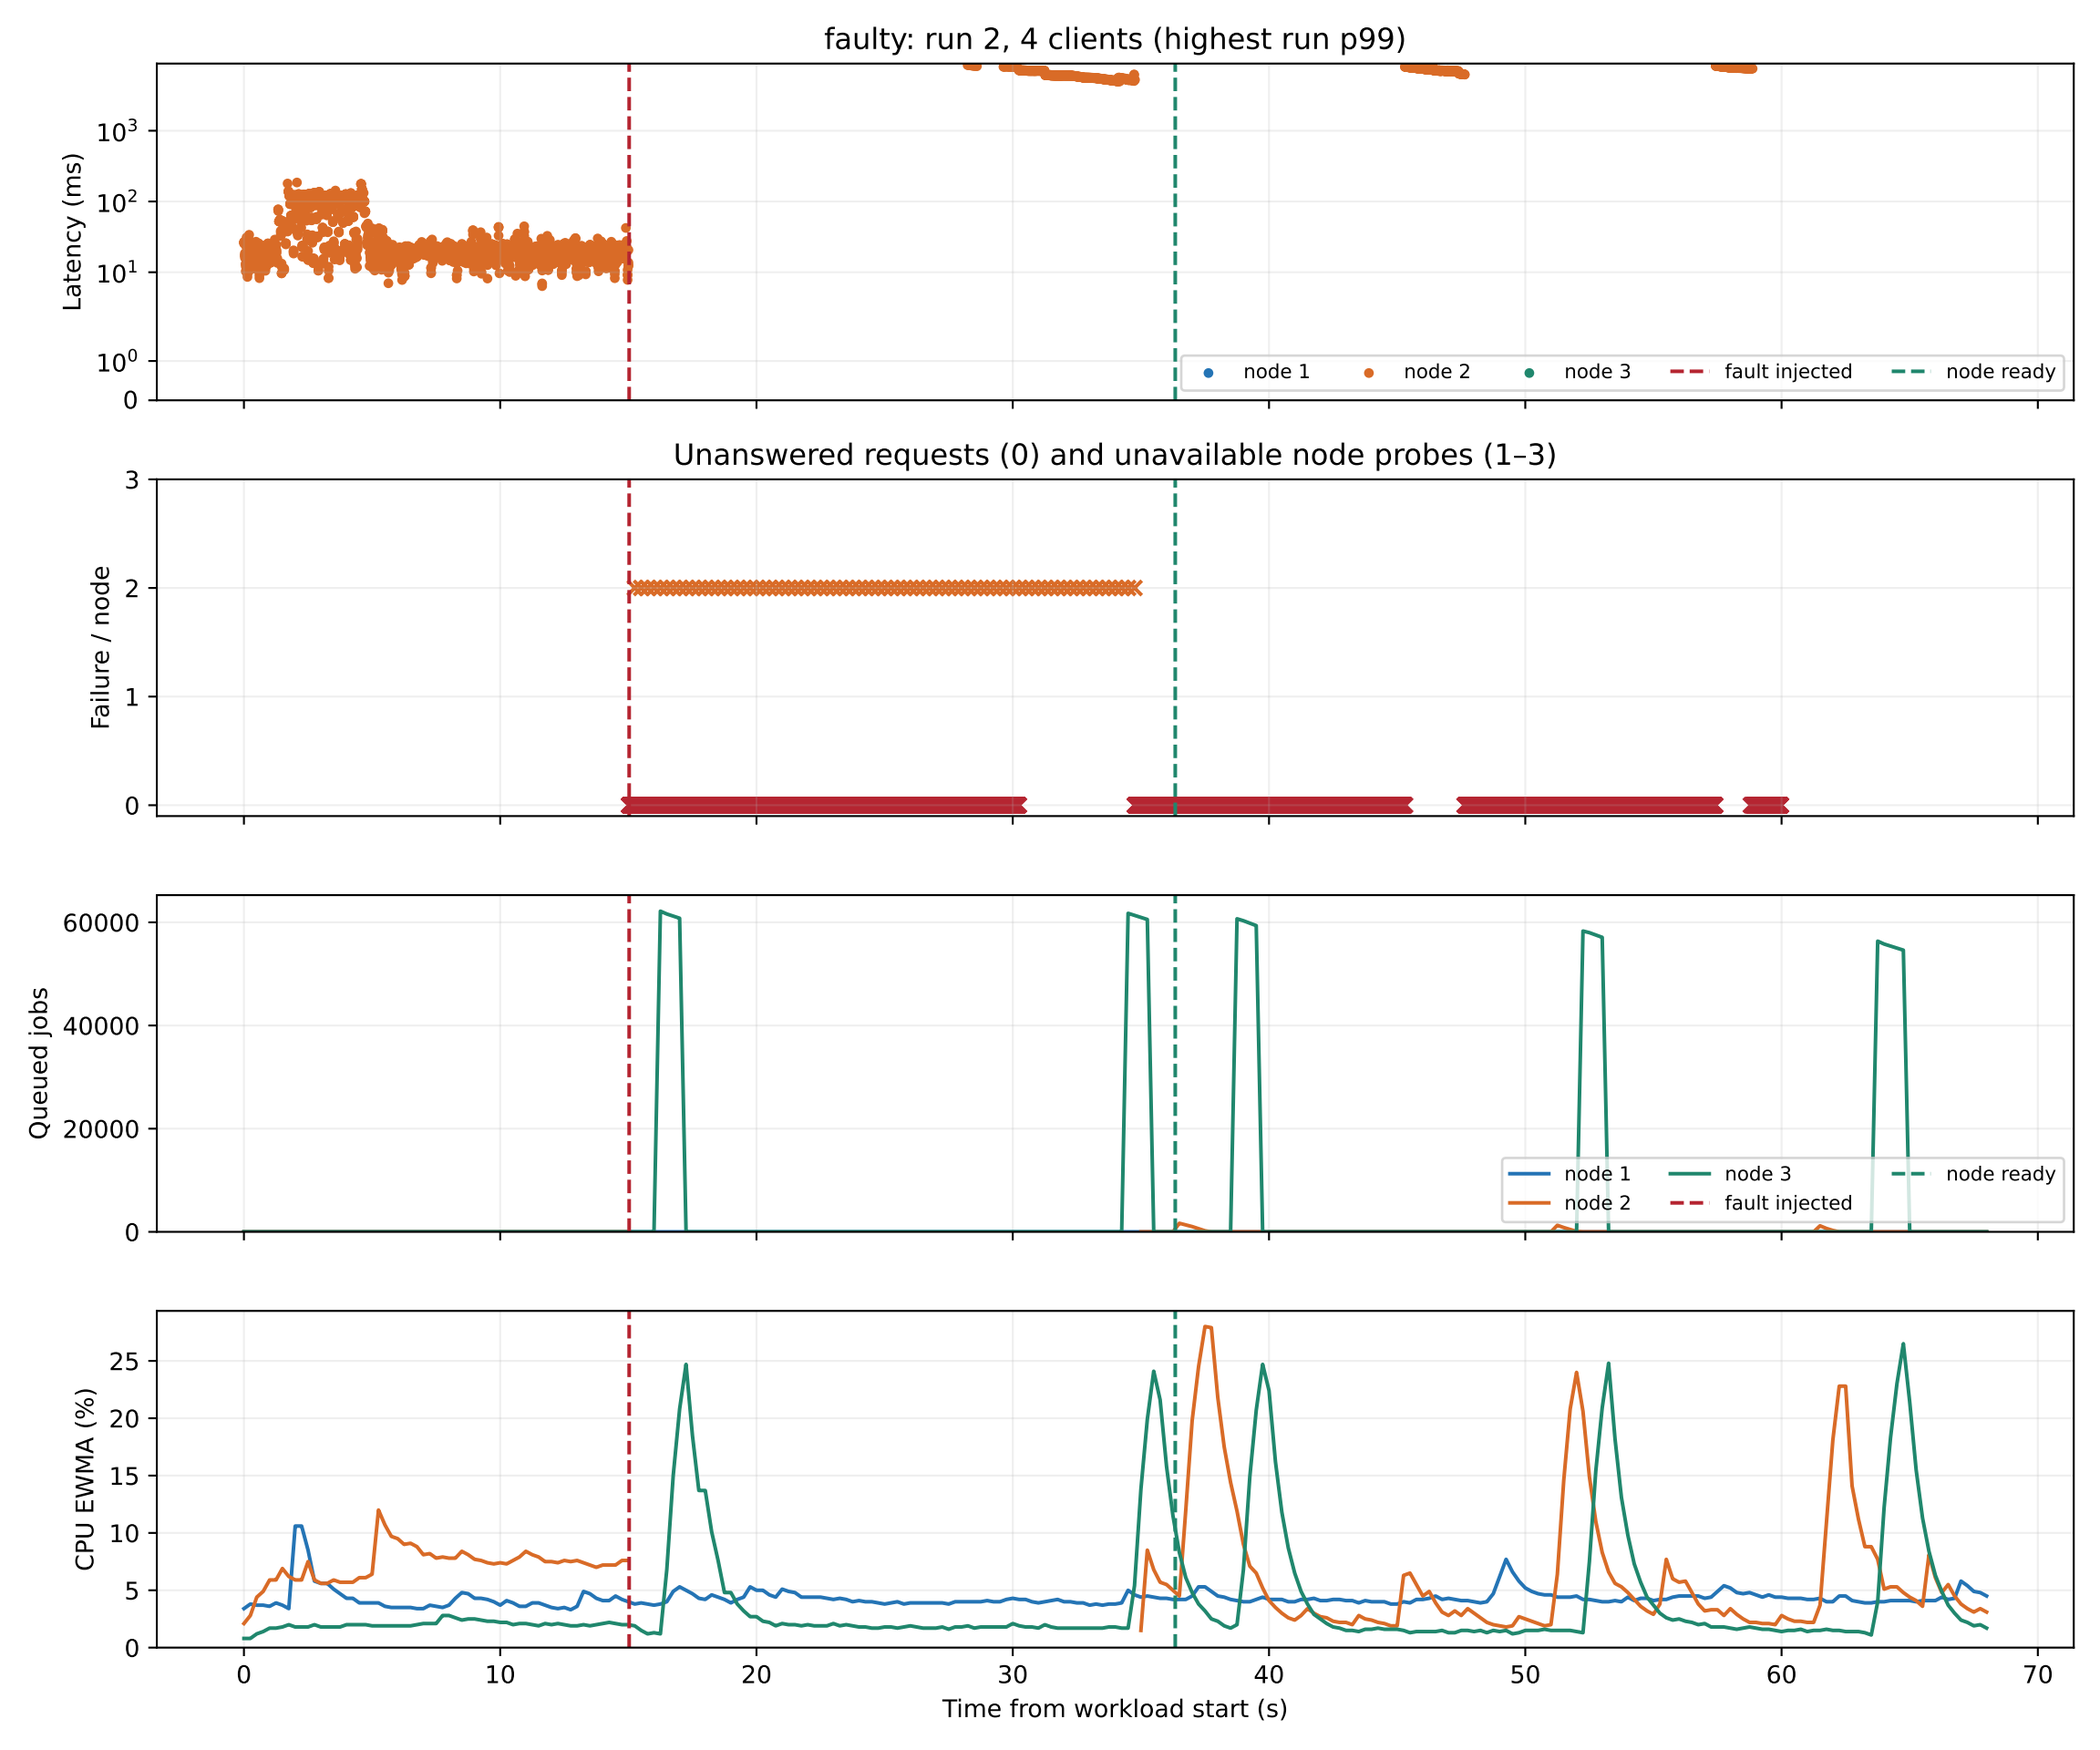
faulty: run 2, 4 clients (highest run p99)
Latency (ms)
node 1
node 2
node 3
fault injected
node ready
Unanswered requests (0) and unavailable node probes (1–3)
Failure / node
Queued jobs
CPU EWMA (%)
Time from workload start (s)